

STAT312: Data Wrangling and Model Selection

01

Fakade, Ncumisa (225091410)

25 July 2025

1 Introduction

Welcome to your first practical on data wrangling in R! This session will teach you the basics of using the `dplyr` package to clean and manipulate data. Data wrangling is a key part of data analysis, and you'll spend a lot of time doing it.

Real data is often messy! It has missing values, inconsistencies, and other problems. We'll learn how to fix these step by step using simple functions from `dplyr`.

The functions we'll learn work together using the pipe operator (`%>%`), which makes your code easier to read.

The Pipe Operator

The pipe operator (`%>%`) allows you to chain functions together, passing the result of one function as the first argument to the next. This creates readable, left-to-right workflows. For example, `data %>% filter(age > 30) %>% select(name)` is equivalent to `select(filter(data, age > 30), name)`. You can access help with `?magrittr::%>%`.

2 Objectives

By the end of this practical, you will be able to:

1. Use basic dplyr functions to manipulate data
2. Clean common data issues
3. Create simple workflows to transform data
4. Practice these skills through exercises

Warning

This is a **long** practical! It will take most students **at least 180 minutes**.

Failure to submit a substantially complete practical will negatively affect your participation mark.

3 Required Packages

```
# Load the tidyverse package
library(tidyverse) # This includes dplyr and other helpful tools
library(lubridate) # For handling dates
library(knitr)      # For making tables
library(rsample)    # For cross-validation functions
library(broom)      # For model output formatting
```

4 Introduction

This combined practical covers the complete workflow from data cleaning to model selection.

5 Dataset Overview

We'll work with two datasets: - hospital_data_messy.csv: 210 patient records with common data quality issues - retail_sales_data.csv: 200 store locations for sales prediction

6 Part 1: Data Wrangling

```
# Import the dataset
hospital_data <- read_csv("hospital_data_messy.csv", na = c("NA", ""))
# Look at the structure
glimpse(hospital_data)

## Rows: 210
## Columns: 8
## $ patient_id      <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14
## $ age              <dbl> 59, 79, 59, 57, 84, 84, 22, 72, 75, 42, NA, 8
## $ gender            <chr> "Male", "Female", "Male", NA, "female", "Male"
## $ admission_date   <chr> "07/01/2023", "02/02/2023", "22/11/2023", "17
## $ department         <chr> "ER", "Surgery", "ER", "Neurology", "emergenc
## $ length_of_stay    <dbl> 3, 5, 11, 3, 7, 9, 5, 5, 8, 3, 4, 7, 6, 9, 1,
## $ treatment_cost     <dbl> 20900, 33700, 11300, 1400, 14500, 9100, 27400
## $ discharge_status  <chr> "Discharged", "Discharged", "Discharged", "Di
```

The dataset has 210 rows (patients) and 8 columns (variables).

```
# Show first few rows  
head(hospital_data)  
  
## # A tibble: 6 x 8  
##   patient_id    age gender admission_date department length_of_stay
```

```

## <dbl> <dbl> <chr> <chr> <chr> <dbl>
## 1 1 59 Male 07/01/2023 ER 3
## 2 2 79 Female 02/02/2023 Surgery 5
## 3 3 59 Male 22/11/2023 ER 11
## 4 4 <NA> 17/04/2023 Neurology 3
## 5 5 female 22/03/2023 emergency 7
## 6 6 84 Male 30/10/2023 A&E 9
## # i 2 more variables: treatment_cost <dbl>, discharge_status <chr>

# Check missing values
hospital_data %>%
  summarise(across(everything(), ~sum(is.na(.)))) %>%
  pivot_longer(
    everything(),
    names_to = "variable",
    values_to = "missing_count"
  ) %>%
  filter(missing_count > 0) %>%
  kable( # kable allows us to make ``nice looking'' tables.
    caption = "Missing data summary",
    format = "pandoc"
)

```

Table 1: Missing data summary

variable	missing_count
patient_id	17
age	14
gender	22
admission_date	21
department	14

variable	missing_count
length_of_stay	14
treatment_cost	12
discharge_status	16

The across Function

The `across()` function applies a function to multiple columns simultaneously. Here, `across(everything(), ~sum(is.na(.)))` calculates the number of missing values for every column. The `~` creates an anonymous function, and `.` represents each column being processed. Check `?dplyr::across` for more details.

The everything Helper

`everything()` is a selection helper that selects all columns in a dataset. Other useful helpers include `starts_with()`, `ends_with()`, `contains()`, and `where()`. These make column selection more flexible and readable. See `?tidyselect::everything` for the complete list.

Reshaping Data with pivot longer

`pivot_longer()` transforms data from wide to long format. Here, it converts our summary of missing values from separate columns into rows with variable names and counts. This format is often more suitable for analysis and visualisation. The reverse operation is `pivot_wider()`. Access help with `?tidyr::pivot_longer`.

7 Step 1: Selecting Columns (Learning objective: `select()` function)

The `select()` function lets you choose which columns to keep. This helps focus on the data you need.

First, we'll trim any extra spaces in text columns.

```
hospital_clean <- hospital_data %>%
  mutate(across(where(is.character), str_trim))
```

The `where` Helper Function

`where()` selects columns based on a logical condition. Here, `where(is.character)` selects all character (text) columns. You can use any function that returns TRUE or FALSE, such as `where(is.numeric)` or `where(is.factor)`. This provides powerful, condition-based column selection. See `?tidyselect::where` for details.

String Trimming with `str_trim`

`str_trim()` removes leading and trailing whitespace from character strings. This is essential for data cleaning as extra spaces can cause matching problems later. The `stringr` package (part of tidyverse) provides many useful string manipulation functions. Check `?stringr::str_trim` for more information.

Example: Select basic patient info.

```
basic_info <- hospital_clean %>%
  select(patient_id, age, gender)

glimpse(basic_info)

## Rows: 210
```

```
## Columns: 3
## $ patient_id <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15,
## $ age         <dbl> 59, 79, 59, 57, 84, 84, 22, 72, 75, 42, NA, 81, 77,
## $ gender      <chr> "Male", "Female", "Male", NA, "female", "Male", "Ma
```

Exercise 1.1: Create a dataset with patient_id, department, and treatment_cost. Call it cost_info.

```
# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code
cost_info <- hospital_clean %>%
  select(patient_id, department, treatment_cost) # Fill in the blanks

# Check
head(cost_info)
```

Exercise 1.2: Select all columns from age to department (using :). Call it demo_data.

```
# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code
demo_data <- hospital_clean %>%
  select(age:department)

# Check
head(demo_data)
```

Column Range Selection

The colon operator (:) selects a range of consecutive columns. For example, age:department selects all columns from age to department inclusive, based on their position in the dataset. This is particularly useful when you need several adjacent columns. You can also use negative selection with - to exclude columns.

8 Step 2: Filtering Rows (Learning objective: filter() function)

The filter() function picks rows that meet conditions.

Example: Patients over 65.

```
elderly_patients <- hospital_clean %>%
  filter(age > 65)

nrow(elderly_patients)

## [1] 61

ncol(elderly_patients)

## [1] 8

hospital_clean%>%
  summarise(across(everything(), ~sum(is.na(.))))%>%
  kable(
    caption= "missing"
)
```

Data Check: Look for problems like impossible ages.

Table 2: missing

patient_id	age	gender	admission_date	department	length_of_stay	treatment_cost	dis
17	14	22	21	14	14	12	

```
suspicious_ages <- hospital_clean %>%
  filter(age < 0 | age > 120)
```

```
negative_stay <- hospital_clean %>%
  filter(length_of_stay < 0)
```

```
negative_cost <- hospital_clean %>%
  filter(treatment_cost < 0)
```

```
zero_cost <- hospital_clean %>%
  filter(treatment_cost == 0)
```

```
# Show if any
suspicious_ages
```

```
## # A tibble: 3 x 8
##   patient_id    age gender admission_date department      length_o
##       <dbl>   <dbl> <chr>     <chr>           <chr>
## 1        148    150 male     30/04/2023 neurology
## 2        189     -5 Male    02/04/2023 Accident & Emergency
## 3        190     999 Male    26/04/2023 Cardiology
```

```
## # i 2 more variables: treatment_cost <dbl>, discharge_status <chr>
```

```
negative_stay
```

```
## # A tibble: 2 x 8
```

```
##   patient_id    age gender admission_date department length_of_stay
```

```

## <dbl> <dbl> <chr> <chr> <chr> <dbl>
## 1 85 27 Male 19/06/2023 Neurology -1
## 2 125 58 MALE 17/09/2023 SURGERY -2
## # i 2 more variables: treatment_cost <dbl>, discharge_status <chr>
negative_cost

## # A tibble: 2 x 8
##   patient_id age gender admission_date department length_of_stay
##       <dbl> <dbl> <chr> <chr> <chr> <dbl>
## 1 48     61 Female 23/01/2023 <NA>      NA
## 2 140    22 Male   05/05/2023 A&E       5
## # i 2 more variables: treatment_cost <dbl>, discharge_status <chr>
nrow(zero_cost)

## [1] 8

```

Logical Operators in filter

The `filter()` function uses logical operators: `>` (greater than), `<` (less than), `==` (equals), `!=` (not equals), `|` (OR), and `&` (AND). You can combine multiple conditions to create complex filters. Missing values require special handling with `is.na()`. See `?base::Logic` for complete operator details.

Exercise 2.1: Filter patients in “Surgery” department. Call it `surgery_patients`. How many are there?

```

# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code
surgery_patients <- hospital_clean %>%
  filter(department == "Surgery")

cat("Number of surgery patients:", nrow(surgery_patients))

```

```
## Number of surgery patients: 13
```

Exercise 2.2: Filter patients under 30 with length_of_stay > 5. Call it young_long_stay.

```
# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code
young_long_stay <- hospital_clean %>%
  filter(age < 30 & length_of_stay > 5)

head(young_long_stay)
```

```
## # A tibble: 5 x 8
##   patient_id    age gender admission_date department      length_o
##       <dbl>   <dbl>  <chr>   <chr>           <chr>
## 1         57     21 Female  18/10/2023  A&E
## 2         NA     19 Male   07/12/2023  Ortho
## 3        149     29 Male   28/03/2023 Neurology
## 4         NA     29 Male   02/04/2023 Cardiology
## 5        189     -5 Male   16/05/2023 Accident & Emergency
## # i 2 more variables: treatment_cost <dbl>, discharge_status <chr>
```

9 Step 3: Transforming Data (Learning objective: mutate() function)

The `mutate()` function adds or changes columns.

Example: Clean age and add age group.

```
hospital_transformed <- hospital_clean %>%
  mutate(
    age_clean = ifelse(age < 0 | age > 120, NA, age),
    age_group = case_when(
```

```

    age_clean < 18 ~ "Child",
    age_clean >= 18 & age_clean < 65 ~ "Adult",
    age_clean >= 65 ~ "Elderly",
    TRUE ~ "Unknown"
)
)

hospital_transformed %>%
  select(patient_id, age, age_clean, age_group) %>%
  head()

## # A tibble: 6 x 4
##   patient_id     age age_clean age_group
##       <dbl>    <dbl>      <dbl>    <chr>
## 1         1      59        59  Adult
## 2         2      79        79  Elderly
## 3         3      59        59  Adult
## 4         4      57        57  Adult
## 5         5      84        84  Elderly
## 6         6      84        84  Elderly

```

Conditional Logic with case_when

`case_when()` provides a powerful way to create multiple conditional statements. It evaluates conditions in order and assigns the first matching result. The `TRUE ~` statement acts as a catch-all for any remaining cases. This is more readable than nested `ifelse()` statements for complex conditions. Check `?dplyr::case_when` for syntax details.

Handling Dates: Dates may have different formats. Use `parse_date_time()` from `lubridate` to standardise them.

First, see some dates.

```
head(hospital_transformed$admission_date)

## [1] "07/01/2023" "02/02/2023" "22/11/2023" "17/04/2023" "22/03/2023"
## [6] "30/10/2023"
```

Now, parse them.

```
hospital_with_dates <- hospital_transformed %>%
  mutate(
    admission_date_clean = parse_date_time(
      admission_date, orders = c("dmy", "mdy", "ymd")))

hospital_with_dates %>%
  select(patient_id, admission_date, admission_date_clean) %>%
  head()
```

```
## # A tibble: 6 x 3
##   patient_id admission_date admission_date_clean
##       <dbl> <chr>           <dttm>
## 1 1 07/01/2023 2023-01-07 00:00:00
## 2 2 02/02/2023 2023-02-02 00:00:00
## 3 3 22/11/2023 2023-11-22 00:00:00
## 4 4 17/04/2023 2023-04-17 00:00:00
## 5 5 22/03/2023 2023-03-22 00:00:00
## 6 6 30/10/2023 2023-10-30 00:00:00
```

Date Parsing with parse_date_time

`parse_date_time()` from the `lubridate` package handles multiple date formats automatically. The `orders` parameter specifies possible formats to try: “dmy” (day-month-year), “mdy” (month-day-year), and “ymd” (year-month-day). The function attempts each format until one succeeds, making it perfect for messy real-world data. See `?lubridate::parse_date_time` for format codes.

Cleaning Text: Standardise department names using simple matches.

First, see variations.

```
head(hospital_with_dates)

## # A tibble: 6 x 11
##   patient_id    age gender admission_date department length_of_stay
##       <dbl>  <dbl> <chr>   <chr>           <chr>          <dbl>
## 1         1     59 Male    07/01/2023    ER              3
## 2         2     79 Female  02/02/2023  Surgery           5
## 3         3     59 Male    22/11/2023    ER             11
## 4         4     57 <NA>    17/04/2023  Neurology        3
## 5         5     84 female  22/03/2023 emergency        7
## 6         6     84 Male    30/10/2023  A&E             9
## # i 5 more variables: treatment_cost <dbl>, discharge_status <chr>,
## #   age_clean <dbl>, age_group <chr>, admission_date_clean <dttm>

hospital_with_dates %>%
  count(department)

## # A tibble: 24 x 2
##   department      n
##   <chr>        <int>
## 1 A&E            22
```

```
## 2 ACCIDENT & EMERGENCY      1
## 3 Accident & Emergency      7
## 4 Bone & Joint              10
## 5 CARDIOLOGY                 7
## 6 Cardiology                  37
## 7 ER                          20
## 8 Emergency                   17
## 9 NEUROLOGY                   2
## 10 Neurology                  24
## # i 14 more rows
```

```
hospital_with_dates %>%
  count(department)
```

```
## # A tibble: 24 x 2
##   department          n
##   <chr>            <int>
## 1 A&E                22
## 2 ACCIDENT & EMERGENCY     1
## 3 Accident & Emergency      7
## 4 Bone & Joint              10
## 5 CARDIOLOGY                 7
## 6 Cardiology                  37
## 7 ER                          20
## 8 Emergency                   17
## 9 NEUROLOGY                   2
## 10 Neurology                  24
## # i 14 more rows
```

```
# Collect together all the different names associated
# with a department
names_for_emergency <- c(
```

```

"emergency", "a&e", "er", "accident & emergency")
names_for_orthopedics <- c(
  "orthopedics", "orthopaedics", "ortho", "bone & joint")
names_for_cardiology <- c("cardiology")
names_for_neurology <- c("neurology")
names_for_surgery <- c("surgery")

# Standardise the department names
hospital_standardised <- hospital_with_dates %>%
  mutate(
    dept_lower = tolower(department),
    department_clean = case_when(
      dept_lower %in% names_for_emergency ~ "Emergency",
      dept_lower %in% names_for_orthopedics ~ "Orthopedics",
      dept_lower %in% names_for_cardiology ~ "Cardiology",
      dept_lower %in% names_for_neurology ~ "Neurology",
      dept_lower %in% names_for_surgery ~ "Surgery",
      TRUE ~ "Unknown"
    ),
    gender = str_to_title(tolower(gender)),
    discharge_status = str_to_title(tolower(discharge_status))
  ) %>%
  select(-dept_lower)

hospital_standardised %>%
  count(department_clean) ##gives a neat standardized count of the departments

## # A tibble: 6 x 2
##   department_clean     n
##   <chr>             <int>
## 1 Emergency           179
## 2 Orthopedics          60
## 3 Cardiology           20
## 4 Neurology            10
## 5 Surgery              10
## 6 Unknown              10

```

## 1 Cardiology	49
## 2 Emergency	71
## 3 Neurology	29
## 4 Orthopedics	32
## 5 Surgery	15
## 6 Unknown	14

The in Operator

The `%in%` operator tests whether elements on the left appear anywhere in the vector on the right. For example, `dept_lower %in% names_for_emergency` returns TRUE if the department name matches any value in the emergency names vector. This is much more efficient than multiple `==` conditions connected with `|`. See `?base::%in%` for details.

String Case Conversion

Functions like `tolower()`, `toupper()`, and `str_to_title()` standardise text case for consistent matching. `str_to_title()` converts text to title case (First Letter Capitalised). Case standardisation is crucial for data cleaning as “Emergency”, “EMERGENCY”, and “emergency” should be treated as identical. The `stringr` package provides additional case conversion options.

Exercise 3.1: Add a column `daily_cost = treatment_cost / length_of_stay`. Call the dataset `hospital_with_daily`.

```
# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.

# Replace underscores (_) with the appropriate code
hospital_with_daily <- hospital_standardised %>%
  mutate(daily_cost = treatment_cost / length_of_stay)
head(hospital_with_daily) ##with all the columns
```

```
hospital_with_daily %>%
  select(patient_id, treatment_cost, length_of_stay, daily_cost) %>% ##with
  head()
```

Exercise 3.2: Add a column high_cost = ifelse(treatment_cost > 20000, "Yes", "No").

```
# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code
hospital_with_high <- hospital_with_daily %>%
  mutate(high_cost = ifelse(treatment_cost > 20000, "Yes", "No"))

hospital_with_high %>%
  select(patient_id, treatment_cost, high_cost) %>%
  head()
```

10 Step 4: Sorting Data (Learning objective: arrange() function)

The arrange() function sorts rows.

Example: Sort by age descending.

```
by_age_desc <- hospital_standardised %>%
  arrange(desc(age_clean))

head(by_age_desc %>% select(patient_id, age_clean))

## # A tibble: 6 x 2
##   patient_id age_clean
##       <dbl>      <dbl>
```

```

## 1      90      85
## 2     106      85
## 3      5      84
## 4      6      84
## 5    138      84
## 6     18      83

```

Sorting with arrange

`arrange()` sorts rows based on one or more columns. Use `desc()` for descending order (largest to smallest). Multiple columns create hierarchical sorting: `arrange(department, desc(age))` sorts by department first, then by age within each department. Missing values typically appear last. See `?dplyr::arrange` for advanced sorting options.

Exercise 4.1: Sort by treatment_cost descending. Call it by_cost_desc.

```

# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code
by_cost_desc <- hospital_standardised %>%
  arrange(desc(treatment_cost))

by_cost_desc %>%
  select(patient_id, treatment_cost) %>%
  head() ##Option A

head(by_cost_desc %>%select(patient_id, treatment_cost)) ##Option B

```

Exercise 4.2: Sort by department_clean ascending, then length_of_stay descending.

```

# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code
by_dept_stay <- hospital_standardised %>%
  arrange(department_clean, desc(length_of_stay))

by_dept_stay %>%
  select(patient_id, department_clean, length_of_stay) %>%
  head()

```

11 Step 5: Summarising Data (Learning objective: summarise() function)

The `summarise()` function calculates summaries, often with `group_by()`.

Example: Overall average age.

```

overall_avg_age <- hospital_standardised %>%
  summarise(avg_age = mean(age_clean, na.rm = TRUE))

overall_avg_age

## # A tibble: 1 x 1
##   avg_age
##       <dbl>
## 1      54.0

```

Grouped: Average cost by department.

```

dept_summary <- hospital_standardised %>%
  group_by(department_clean) %>%
  summarise(avg_cost = mean(treatment_cost, na.rm = TRUE))

```

```
dept_summary
```

```
## # A tibble: 6 x 2
##   department_clean avg_cost
##   <chr>              <dbl>
## 1 Cardiology          26962.
## 2 Emergency           16055.
## 3 Neurology            12750
## 4 Orthopedics          15700
## 5 Surgery              39160
## 6 Unknown              21485.
```

Grouped Operations

`group_by()` creates groups within your data, causing subsequent operations to be performed separately for each group. Here, `group_by(department_clean)` means the `summarise()` function calculates separate averages for each department. This is fundamental for comparative analysis. Always check your grouping with `groups()` and remove with `ungroup()` when finished. See `?dplyr::group_by`.

Handling Missing Values in Summaries

The `na.rm = TRUE` argument tells summary functions like `mean()`, `sum()`, and `sd()` to ignore missing values. Without this, any missing value causes the entire summary to return `NA`. This parameter is essential when working with real-world data containing missing observations. Check each function's help page for missing value handling options.

Exercise 5.1: Calculate total patients and average stay overall. Call it `overall_summary`.

```
hospital_standardised
```

```
## # A tibble: 210 x 12
```

```

##      patient_id    age gender admission_date department    length_of_stay
##              <dbl> <dbl> <chr>   <chr>          <chr>           <dbl>
## 1            1     59 Male    07/01/2023    ER             3
## 2            2     79 Female  02/02/2023  Surgery          5
## 3            3     59 Male    22/11/2023    ER            11
## 4            4     57 <NA>    17/04/2023 Neurology         3
## 5            5     84 Female  22/03/2023 emergency        7
## 6            6     84 Male    30/10/2023 A&E            9
## 7            7     22 Male    28/10/2023 Orthopaedics      5
## 8            8     72 Male    01/11/2023 CARDIOLOGY       5
## 9            9     75 Female  07/12/2023    ER            8
## 10           10    42 Female  01/02/2023    ER            3
## # i 200 more rows
## # i 6 more variables: treatment_cost <dbl>, discharge_status <chr>,
## #   age_clean <dbl>, age_group <chr>, admission_date_clean <dttm>,
## #   department_clean <chr>

# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code

overall_summary <- hospital_standardised %>%
  summarise(
    total_patients = n(),
    avg_stay = mean(length_of_stay, na.rm = TRUE)
  )

overall_summary

```

The n Function

`n()` counts the number of rows in each group. When used without `group_by()`, it counts all rows. When used with grouping, it provides counts for each group separately. This is equivalent to `nrow()` for ungrouped data but works properly with grouped operations. Other useful counting functions include `n_distinct()` for unique values. See `?dplyr::n`.

Exercise 5.2: Group by gender, summarise average age and count.

```
# NB: After editing the code below, change the code chunk
# option eval=FALSE to eval=TRUE.
# Replace underscores (_) with the appropriate code
gender_summary <- hospital_standardised %>%
  group_by(gender) %>%
  summarise(
    avg_age = mean(age, na.rm = TRUE),
    count = n()
  )

gender_summary
```

12 Step 6: Putting It Together with Pipes

Use `%>%` to chain steps.

Example: Filter elderly, summarise average cost.

```
elderly_avg_cost <- hospital_standardised %>%
  filter(age_group == "Elderly") %>%
  summarise(avg_cost = mean(treatment_cost, na.rm = TRUE))

elderly_avg_cost
```

```

## # A tibble: 1 x 1
##   avg_cost
##   <dbl>
## 1 20885.

```

Exercise 6: Filter “Emergency” department, sort by descending treatment_cost, select top 5.

hospital_standardised

```

## # A tibble: 210 x 12
##   patient_id    age gender admission_date department length_of_stay
##   <dbl> <dbl> <chr>   <chr>           <chr>          <dbl>
## 1 1       59  Male    07/01/2023  ER             3
## 2 2       79  Female  02/02/2023  Surgery        5
## 3 3       59  Male    22/11/2023  ER             11
## 4 4       57  <NA>    17/04/2023  Neurology      3
## 5 5       84  Female  22/03/2023  emergency     7
## 6 6       84  Male    30/10/2023  A&E           9
## 7 7       22  Male    28/10/2023  Orthopaedics  5
## 8 8       72  Male    01/11/2023  CARDIOLOGY   5
## 9 9       75  Female  07/12/2023  ER             8
## 10 10      42  Female  01/02/2023  ER             3
## # i 200 more rows
## # i 6 more variables: treatment_cost <dbl>, discharge_status <chr>,
## #   age_clean <dbl>, age_group <chr>, admission_date_clean <dttm>,
## #   department_clean <chr>

```

NB: After editing the code below, change the code chunk

option eval=FALSE to eval=TRUE.

Replace underscores (_) with the appropriate code

```
top_emergency <- hospital_standardised %>%

```

```
  filter(department== "Emergency") %>%
```

```
arrange(desc(treatment_cost)) %>%  
  slice(1:5) # top 5  
  
top_emergency
```

Row Selection with slice

`slice()` selects rows by their position. `slice(1:5)` selects the first five rows, whilst `slice(c(1, 3, 5))` selects specific positions. Related functions include `slice_head()`, `slice_tail()`, `slice_max()`, and `slice_min()` for more targeted selection. This is particularly useful after sorting to get top or bottom values. See `?dplyr::slice` for all variants.

13 Step 7: Final Clean Dataset

Here's a full cleaning pipeline.

```
# Collect together all the different names associated  
# with a department  
names_for_emergency <- c("emergency", "a&e", "er",  
                           "accident & emergency")  
names_for_orthopedics <- c("orthopedics", "orthopaedics",  
                            "ortho", "bone & joint")  
names_for_cardiology <- c("cardiology")  
names_for_neurology <- c("neurology")  
names_for_surgery <- c("surgery")  
  
hospital_final <- hospital_data %>%  
  mutate(across(where(is.character), str_trim)) %>%  
  mutate(  
    age_clean = ifelse(age < 0 | age > 120, NA, age),  
    length_of_stay = ifelse(length_of_stay < 0, NA, length_of_stay),
```

```

treatment_cost = ifelse(treatment_cost < 0, NA, treatment_cost),
admission_date_clean = parse_date_time(
  admission_date, orders = c("dmy", "mdy", "ymd")),
dept_lower = tolower(department),
department_clean = case_when(
  dept_lower %in% names_for_emergency ~ "Emergency",
  dept_lower %in% names_for_orthopedics ~ "Orthopedics",
  dept_lower %in% names_for_cardiology ~ "Cardiology",
  dept_lower %in% names_for_neurology ~ "Neurology",
  dept_lower %in% names_for_surgery ~ "Surgery",
  TRUE ~ "Unknown"
),
gender = str_to_title(tolower(gender)),
discharge_status = str_to_title(tolower(discharge_status)),
age_group = case_when(
  age_clean < 18 ~ "Child",
  age_clean >= 18 & age_clean < 65 ~ "Adult",
  age_clean >= 65 ~ "Elderly",
  TRUE ~ "Unknown"
),
daily_cost = treatment_cost / length_of_stay
) %>%
select(-dept_lower) %>%
distinct() %>%
select(patient_id, age_clean, gender, admission_date_clean,
       department_clean, length_of_stay, treatment_cost,
       discharge_status, age_group, daily_cost)

head(hospital_final)

```

```
## # A tibble: 6 x 10
##   patient_id age_clean gender admission_date_clean department_clean
##       <dbl>      <dbl>   <chr>        <dttm>                  <chr>
## 1           1        59 Male  2023-01-07 00:00:00 Emergency
## 2           2        79 Female 2023-02-02 00:00:00 Surgery
## 3           3        59 Male  2023-11-22 00:00:00 Emergency
## 4           4        57 <NA>  2023-04-17 00:00:00 Neurology
## 5           5        84 Female 2023-03-22 00:00:00 Emergency
## 6           6        84 Male  2023-10-30 00:00:00 Emergency
## # i 5 more variables: length_of_stay <dbl>, treatment_cost <dbl>,
## #   discharge_status <chr>, age_group <chr>, daily_cost <dbl>
```

Removing Duplicates with distinct

`distinct()` removes duplicate rows from your dataset. By default, it considers all columns when identifying duplicates. You can specify particular columns with `distinct(column1, column2)` to remove duplicates based only on those variables. Use `.keep_all = TRUE` to retain all columns when specifying particular variables. This is essential for data quality assurance. See `?dplyr::distinct` for options.

Export:

```
write_csv(hospital_final, "hospital_data_clean.csv")
```

14 Transition: From Clean Data to Business Analytics

The cleaning skills you've just mastered aren't just academic—they're essential for real business problems. The retail data we'll now analyze has similar issues: missing values, negative entries, and outliers that need handling before modeling.

Notice how we'll use the same `mutate()`, `ifelse()`, and `filter()` functions to prepare this data for analysis.

15 Part 2: Model Selection

```
# Clear workspace from Part 1  
rm(list = ls()[!ls() %in% c("hospital_data")]) # Keep only if needed
```

16 Business Context

RetailMax is a chain of electronics stores. They want to improve monthly sales performance across their 200 store locations. The analytics team has collected data on various factors that might influence sales. These include advertising spending, staffing levels, store features, and local market conditions.

Your task is to develop a model to predict monthly sales revenue. This will help management make better decisions about resource allocation and planning.

17 Dataset Import and Exploration

```
# Import the sales data  
sales_data <- read_csv("retail_sales_data.csv")  
  
# Initial exploration  
glimpse(sales_data)  
  
## Rows: 200  
## Columns: 8  
## $ monthly_sales      <dbl> -150000, 678303, 755785, 831246, 854256,  
## $ advertising_spend   <dbl> 14599, 71756, 13050, 60465, 49036, 12695,
```

```

## $ staff_count <dbl> 27, 11, 16, 5, 16, 24, 16, 14, 13, 15, 22
## $ store_size <dbl> 1044, 224, 1122, 426, 150, 858, 1192, 896
## $ foot_traffic <dbl> NA, 8477, 6055, 6837, 5725, 4243, NA, 5370
## $ local_income <dbl> 32798, 12000, 17894, 17044, 29601, 12000,
## $ competition_distance <dbl> 5.5, 0.4, 5.4, 3.3, 6.0, 6.8, 2.9, 0.4, 8
## $ seasonal_index <dbl> 1.1998192, 0.9874935, 1.1563594, 0.700381

```

The dataset contains 200 monthly observations across 8 variables:

- **monthly_sales**: Total sales revenue (R, target variable)
- **advertising_spend**: Monthly advertising spending (R)
- **staff_count**: Number of full-time equivalent employees
- **store_size**: Retail floor space (square metres)
- **foot_traffic**: Monthly customer visits
- **local_income**: Average household income in area (R)
- **competition_distance**: Distance to nearest competitor (km)
- **seasonal_index**: Seasonal adjustment factor (1.0 = average)

```
# Examine the first few rows
```

```
head(sales_data)
```

```

## # A tibble: 6 x 8
##   monthly_sales advertising_spend staff_count store_size foot_traffic
##   <dbl>           <dbl>        <dbl>       <dbl>          <dbl>
## 1 -150000            14599        27         1044          NA
## 2    678303            71756        11         224         8477
## 3    755785            13050        16        1122         6055
## 4    831246            60465         5         426         6837
## 5    854256            49036        16         150         5725
## 6    924861            12695        24         858         4243
## # i 3 more variables: local_income <dbl>, competition_distance <dbl>,
## #   seasonal_index <dbl>

```

```
# Generate summary statistics  
summary(sales_data)
```

```
## monthly_sales      advertising_spend   staff_count      store_size  
## Min. : -150000    Min. :-25000     Min. : 3.00    Min. : 150.0  
## 1st Qu.: 1254575  1st Qu.: 61169    1st Qu.: 13.75   1st Qu.: 529.2  
## Median : 1452447  Median : 86009    Median : 19.00   Median : 788.5  
## Mean   : 1631742  Mean   : 86870    Mean   : 19.50   Mean   : 855.4  
## 3rd Qu.: 1692400  3rd Qu.:110033   3rd Qu.: 23.25   3rd Qu.:1095.2  
## Max.   :25000000  Max.   :450000    Max.   :150.00   Max.   :8500.0  
##                               NA's :11  
## foot_traffic      local_income      competition_distance seasonal_index  
## Min. : 500        Min. : 12000     Min. : -2.500    Min. : 0.6299  
## 1st Qu.: 4740      1st Qu.: 24652    1st Qu.: 0.875    1st Qu.: 0.9027  
## Median : 6128      Median : 35002    Median : 2.550    Median : 1.0051  
## Mean   : 6957      Mean   : 39797    Mean   : 3.083    Mean   : 1.0160  
## 3rd Qu.: 8306      3rd Qu.: 45322    3rd Qu.: 4.400    3rd Qu.: 1.1054  
## Max.   :85000      Max.   :850000    Max.   :17.900    Max.   :1.3842  
## NA's   :6          NA's   :5
```

Task 1: Find potential data quality problems from the summary statistics. Look for:

- Variables with missing values (NAs)
- Unrealistic values (negative numbers where they should not exist)
- Extreme outliers that seem impossible

Write your findings below:

Data Quality Issues Found

advertising spend has 11 missing values, foot traffic has 6 missing values, local income has 5 missing values, there are unrealistic values, extremely negative

minimuns, there extreme outliers in all columns except for seasonal index

18 Data Cleaning

Before building models, fix the data quality issues you found.

18.1 Missing Values

Hint: Some variables have missing values that need attention. Think about whether missing values represent real missingness or data entry errors. For this business context, simple imputation methods (such as median imputation for continuous variables) work well.

```
#Countmissingvaluespervariable
missing_summary<-sales_data %>%
  summarise_all(~sum(is.na(.))) %>%
  pivot_longer(
    everything(),
    names_to="variables",
    values_to="missing_count"
)%>%
  filter(missing_count>0)
missing_summary
```

```
## # A tibble: 3 x 2
##   variables      missing_count
##   <chr>              <int>
## 1 advertising_spend      11
## 2 foot_traffic            6
## 3 local_income             5
```

```

# Your code here: Find and handle missing values
# Suggestion: Use is.na() to identify missing patterns
# Consider median or mean imputation for continuous variables

# 

#Count missing values per variable
missing_summary<-sales_data %>%
  summarise_all(~sum(is.na(.))) %>%
  gather(variables,missing_count)%>%
  filter(missing_count>0)

#missing_summary

#Identify negative values in variables that should be non-negative
negative_vars<-c("monthly_sales","advertising_spend","staff_count","store
                 "foot_traffic","local_income","competition_distance")

negative_summary<-sales_data %>%
  select(all_of(negative_vars))%>%
  summarise_all(~sum(.<0, na.rm=TRUE )) %>%
  gather(variables,negative_count)%>%
  filter(negative_count>0)

#Calculate extreme outliers using IQR method
outlier_summary<-sales_data %>%
  select_if(is.numeric)%>%
  summarise_all(~ {
    q1<-quantile(.,0.25, na.rm= TRUE)
    q3 <- quantile(., 0.75, na.rm = TRUE)
    iqr <- q3-q1
  })

```

```

    sum(. < (q1-1.5*iqr) | . > (q3+1.5*iqr), na.rm = TRUE)
}) %>%
gather(variables,outlier_count)%>%
filter(outlier_count>0)

```

18.2 Outlier Detection

Hint: Look at variables for extreme values that could be data entry errors or genuinely unusual observations. Pay special attention to values that are negative when they should be positive, or values that are much larger than typical observations.

```

# Your code here: Find and handle outliers
# Suggestion: Use summary() and boxplots to identify extreme values
# Consider the business context when deciding if values are realistic
# Once identified the outliers, use the IQR method to remove them:
#   IQR Method: Values beyond Q1 - 1.5×IQR or Q3 + 1.5×IQR
#               may be statistical outliers

# Store original row count for tracking
original_rows <- nrow(sales_data)

#Handle missing values with median imputation
sales_data<-sales_data %>%
mutate(
  advertising_spend= if_else(
    is.na(advertising_spend),
    median(advertising_spend, na.rm = TRUE),
    advertising_spend),

```

```

foot_traffic= if_else(
  is.na(foot_traffic),
  median(foot_traffic,na.rm = TRUE),
  foot_traffic),

local_income= if_else(
  is.na(local_income),
  median(local_income,na.rm = TRUE),
  local_income)

)
#Verify no missing values remain
missing_after_imputation<-sum(is.na(sales_data))
missing_after_imputation

## [1] 0

#Convert negative values to positive (assuming data entry errors)
sales_data<-sales_data %>%
  mutate(
    monthly_sales= abs(monthly_sales),
    advertising_spend= abs(advertising_spend),
    competition_distance= abs(competition_distance)
  )

#Function to remove outliers using IQR method
remove_outliers<-function(data,var) {
  #Calculate quartiles and IQR
  q1<-quantile(data[[var]],0.25, na.rm= TRUE)
  q3<-quantile(data[[var]],0.75, na.rm= TRUE)
  iqr<-q3-q1
}

```

```

#Define outlier bounds (remeber the sum structure)
lower_bound<-q1-1.5*iqr
upper_bound<-q3+1.5*iqr

#Filter data to exclude outliers
filtered_data<-data %>%
  filter(.data[[var]]>= lower_bound & .data[[var]]<= upper_bound)
return(filtered_data)
}

#Apply outlier removal to key variables
sales_data<-sales_data %>%
  remove_outliers("monthly_sales")%>%
  remove_outliers("advertising_spend")%>%
  remove_outliers("staff_count")%>%
  remove_outliers("store_size")%>%
  remove_outliers("foot_traffic")%>%
  remove_outliers("local_income")%>%
  remove_outliers("competition_distance")

# Calculate cleaning impact#####
cleaned_rows<- nrow(sales_data)
rows_removed <- original_rows- cleaned_rows

cleaned_rows

## [1] 178

```

```
rows_removed
```

```
## [1] 22
```

Task 2: Explain your data cleaning decisions. Why did you choose specific imputation methods and outlier handling strategies?

Data Cleaning Explanation:

I imputed the missing values using the median because the data was skewed and contained outliers. The mean is sensitive to outliers, which could distort the results. For negative values, I applied the absolute value function to convert them to positive values, as the negatives were likely due to human error and it is more appropriate to work with positive figures in this context. I then removed outliers on a per-column basis, as this approach is more precise and ensures targeted data cleaning.

19 Model Specification and Cross-Validation

Build multiple linear regression models with different predictor combinations. Use k-fold cross-validation to compare model performance fairly.

19.1 Model Specifications

Consider these potential model specifications:

1. **Simple Model:** Key predictors only
2. **Full Model:** All available predictors
3. **Optimised Model:** Carefully chosen subset based on business logic

```
# Your code here: Define your model formulas
```

```
# Example structure (modify as needed):
```

```
model_1 <- monthly_sales ~ advertising_spend + staff_count
```

```
model_2 <- monthly_sales ~ . # Full model  
model_3 <- monthly_sales ~ advertising_spend + staff_count + store_size +
```

19.2 Cross-Validation Implementation

Use 10-fold cross-validation to evaluate each model specification.

Hint: Use `rsample::vfold_cv()` to create the cross-validation folds. The `rsample` package works well with `dplyr` workflows using `map()` functions for applying models across folds.

Your code here: Set up cross-validation folds

Suggestion: Use `vfold_cv()` with `v = 10`

```
set.seed(123) #Forreproducibility  
cv_folds<- vfold_cv(sales_data, v=10)
```

Your code here: Apply cross-validation to each model

Suggestion: Create a function that fits a model to training data

and calculates performance metrics on validation data

Use `map()` to apply this function across all folds

```
calc_metrics<-function(split,model_formula) {  
  train <- analysis(split)  
  test <- assessment(split)  
  fit <- lm(model_formula,data = train)  
  preds <- predict(fit,newdata =test)}
```

#Calculate performance metrics#####

```
rmse<-sqrt(mean((test$monthly_sales-preds)^2))
```

```
mae<-mean(abs(test$monthly_sales-preds))
```

```

r2<-1-sum((test$monthly_sales-preds)^2)/sum((test$monthly_sales-mean(test$monthly_sales))^2)

tibble(rmse= rmse, mae = mae, r2 =r2)
}

#Apply cross-validation to each model

results_1 <- cv_folds %>%
  mutate(metrics= map(splits,~calc_metrics(.,model_1))) %>%
  unnest(metrics)

results_2 <- cv_folds %>%
  mutate(metrics= map(splits,~calc_metrics(.,model_2))) %>%
  unnest(metrics)

results_3 <- cv_folds %>%
  mutate(metrics= map(splits,~calc_metrics(.,model_3))) %>%
  unnest(metrics)

```

19.3 Performance Metrics

Calculate and compare relevant performance metrics for each model:

- **RMSE** (Root Mean Square Error): Measures average prediction error
- **MAE** (Mean Absolute Error): Measures average absolute prediction error
- **R²** (R-squared): Proportion of variance explained

```

# Your code here: Calculate performance metrics for each model
# Compare mean performance across all cross-validation folds

```

```

#Summarise metrics with confidence intervals

```

```

summary_1<-results_1 %>%

```

```

summarise(
  rmse_mean = mean(rmse),
  rmse_se = sd(rmse) /sqrt(n()),
  mae_mean =mean(mae),
  mae_se =sd(mae) /sqrt(n()),
  r2_mean =mean(r2),
  r2_se =sd(r2) /sqrt(n())
) %>%
  mutate(model= "Simple")

summary_2<-results_2 %>%
  summarise(
    rmse_mean = mean(rmse),
    rmse_se = sd(rmse) /sqrt(n()),
    mae_mean =mean(mae),
    mae_se =sd(mae) /sqrt(n()),
    r2_mean =mean(r2),
    r2_se =sd(r2) /sqrt(n())
) %>%
  mutate(model= "Full")

summary_3<-results_3 %>%
  summarise(
    rmse_mean = mean(rmse),
    rmse_se = sd(rmse) /sqrt(n()),
    mae_mean =mean(mae),
    mae_se =sd(mae) /sqrt(n()),
    r2_mean =mean(r2),
    r2_se =sd(r2) /sqrt(n())
) %>%

```

```

mutate(model= "Optimised")

#Combine summaries
all_summaries<-bind_rows(summary_1, summary_2, summary_3)

#Identify best performing model
best_model_name<-all_summaries%>%
arrange(rmse_mean)%>%
slice(1)%>%
pull(model)

```

Task 3: Present your cross-validation results in a clear, professional format. Create a summary table showing mean performance metrics and their standard errors across folds.

```

# Your code here: Create a summary table of results

#Create comprehensive summary table
all_summaries %>%
  select(model, rmse_mean, rmse_se, mae_mean, mae_se, r2_mean, r2_se) %>%
  kable(
    digits = c(0, rep(2,6)), # 0 decimal places for the first column ("Model")
    caption = "Cross-Validation Performance Metrics",
    col.names = c("Model", "RMSE (Mean)", "RMSE (SE)",
                 "MAE (Mean)", "MAE (SE)", "R2 (Mean)", "R2 (SE)"),
    format = "pandoc"
  )

```

Table 3: (#tab:results_table) Cross-Validation Performance Metrics

Model	RMSE (Mean)	RMSE (SE)	MAE (Mean)	MAE (SE)	R2 (Mean)	R2 (SE)
Simple	277166.2	17487.45	225521.1	15240.07	0.13	0.04
Full	196198.8	10900.57	151168.3	8403.76	0.55	0.03
Optimised	235012.6	14854.16	191730.3	12927.75	0.36	0.05

20 Model Selection and Interpretation

Based on your cross-validation results, select the best model specification.

Task 4: Explain your model selection decision. Consider both statistical performance and business interpretability.

Model Selection Explanation:

Based on the cross-validation results, the full model was selected because it achieved a low RMSE and MAE, a high R², and demonstrated stable performance across all folds. This balance between accuracy and consistency makes it the most suitable model for business use

20.1 Final Model Analysis

Fit your selected model to the complete dataset and examine the results.

```
# Your code here: Fit the selected model to the full dataset
# Examine coefficient estimates, significance, and overall fit

#Extract performance metrics for best model
best_metrics<-all_summaries%>%
  filter(model== best_model_name )
```

```

#Calculate performance improvements
simple_rmse<-all_summaries%>%
  filter(model=="Simple")%>%
  pull(rmse_mean)

best_rmse<-best_metrics%>%
  pull(rmse_mean)

rmse_improvement<-
  ( simple_rmse-best_rmse)/ simple_rmse) *100

#Compare_all_models
rmse_comparison<-all_summaries%>%
  select(model,rmse_mean, r2_mean) %>%
  arrange(rmse_mean)

# Fit the selected model based on CV results
if(best_model_name == "Simple") {
  final_fit <- lm(model_1, data =sales_data)
  selected_formula<-model_1
} else if(best_model_name == "Full") {
  final_fit <- lm(model_2, data = sales_data)
  selected_formula<-model_2
} else {
  final_fit <- lm(model_3, data = sales_data)
  selected_formula<-model_3
}

# Display model summary
summary(final_fit)

```

```

## Call:
## lm(formula = model_2, data = sales_data)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -527325 -114179    2096 106772 558702
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)
## (Intercept)             -94298.870 132670.278 -0.711 0.47820
## advertising_spend        2.601      0.418   6.222 3.69e-09 ***
## staff_count              15671.981  2102.067   7.456 4.36e-12 ***
## store_size                252.502    38.817   6.505 8.35e-10 ***
## foot_traffic               47.786    5.824   8.205 5.48e-14 ***
## local_income                 8.912    1.029   8.663 3.47e-15 ***
## competition_distance -12011.502   6278.238  -1.913 0.05740 .
## seasonal_index            276401.130  97046.885   2.848 0.00494 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 194500 on 170 degrees of freedom
## Multiple R-squared:  0.6195, Adjusted R-squared:  0.6038
## F-statistic: 39.54 on 7 and 170 DF,  p-value: < 2.2e-16

# Extract key model statistics
model_stats <- glimpse(final_fit)

## List of 12
## $ coefficients : Named num [1:8] -94298.9 2.6 15672 252.5 47.8 ...
## ..- attr(*, "names")= chr [1:8] "(Intercept)" "advertising_spend" "s
```

```

## $ residuals      : Named num [1:178] -423162 -421488 -50214 -203327 -108...
##   ..- attr(*, "names")= chr [1:178] "1" "2" "3" "4" ...
## $ effects       : Named num [1:178] -19740035 903638 1527237 1196025 15...
##   ..- attr(*, "names")= chr [1:178] "(Intercept)" "advertising_spend" ...
## $ rank          : int 8
## $ fitted.values: Named num [1:178] 1101465 1177273 881460 1057583 103...
##   ..- attr(*, "names")= chr [1:178] "1" "2" "3" "4" ...
## $ assign         : int [1:8] 0 1 2 3 4 5 6 7
## $ qr            :List of 5
##   ..$ qr     : num [1:178, 1:8] -13.342 0.075 0.075 0.075 0.075 ...
##   ... ..- attr(*, "dimnames")=List of 2
##   ... ...- attr(*, "assign")= int [1:8] 0 1 2 3 4 5 6 7
##   ..$ qraux: num [1:8] 1.07 1.15 1.13 1.13 1.07 ...
##   ..$ pivot: int [1:8] 1 2 3 4 5 6 7 8
##   ..$ tol   : num 1e-07
##   ..$ rank  : int 8
##   ..- attr(*, "class")= chr "qr"
## $ df.residual   : int 170
## $ xlevels       : Named list()
## $ call          : language lm(formula = model_2, data = sales_data)
## $ terms         :Classes 'terms', 'formula' language monthly_sales ~ ...
##   .. ..- attr(*, "variables")= language list(monthly_sales, advertising_spend)
##   .. ..- attr(*, "factors")= int [1:8, 1:7] 0 1 0 0 0 0 0 0 ...
##   .. ... ..- attr(*, "dimnames")=List of 2
##   .. ... ..- attr(*, "term.labels")= chr [1:7] "advertising_spend" "staff_size" ...
##   .. ... ..- attr(*, "order")= int [1:7] 1 1 1 1 1 1 1
##   .. ... ..- attr(*, "intercept")= int 1
##   .. ... ..- attr(*, "response")= int 1
##   .. ... ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
##   .. ... ..- attr(*, "predvars")= language list(monthly_sales, advertising_spend)

```

```
## ... - attr(*, "dataClasses")= Named chr [1:8] "numeric" "numeric" ...
## ... . . - attr(*, "names")= chr [1:8] "monthly_sales" "advertising_sp...
## $ model      : 'data.frame':   178 obs. of  8 variables:
##   ..$ monthly_sales       : num [1:178] 678303 755785 831246 854256 92...
##   ..$ advertising_spend  : num [1:178] 71756 13050 60465 49036 12695 ...
##   ..$ staff_count         : num [1:178] 11 16 5 16 24 16 14 13 22 19 ...
##   ..$ store_size          : num [1:178] 224 1122 426 150 858 ...
##   ..$ foot_traffic        : num [1:178] 8477 6055 6837 5725 4243 ...
##   ..$ local_income         : num [1:178] 12000 17894 17044 29601 12000 ...
##   ..$ competition_distance: num [1:178] 0.4 5.4 3.3 6 6.8 2.9 0.4 8.8 ...
##   ..$ seasonal_index       : num [1:178] 0.987 1.156 0.7 0.978 0.989 ...
##   ..- attr(*, "terms")=Classes 'terms', 'formula' language monthly_sa...
##   ... . . . - attr(*, "variables")= language list(monthly_sales, adverti...
##   ... . . . - attr(*, "factors")= int [1:8, 1:7] 0 1 0 0 0 0 0 0 ...
##   ... . . . . - attr(*, "dimnames")=List of 2
##   ... . . . . - attr(*, "term.labels")= chr [1:7] "advertising_spend" "sta...
##   ... . . . . - attr(*, "order")= int [1:7] 1 1 1 1 1 1 1
##   ... . . . . - attr(*, "intercept")= int 1
##   ... . . . . - attr(*, "response")= int 1
##   ... . . . . - attr(*, ".Environment")=<environment: R_GlobalEnv>
##   ... . . . . - attr(*, "predvars")= language list(monthly_sales, adverti...
##   ... . . . . - attr(*, "dataClasses")= Named chr [1:8] "numeric" "numeric" ...
##   ... . . . . . - attr(*, "names")= chr [1:8] "monthly_sales" "advertising_...
## - attr(*, "class")= chr "lm"

coef_summary <- glimpse(final_fit)
```

```
## List of 12
## $ coefficients : Named num [1:8] -94298.9 2.6 15672 252.5 47.8 ...
##   ..- attr(*, "names")= chr [1:8] "(Intercept)" "advertising_spend" "s
## $ residuals      : Named num [1:178] -423162 -421488 -50214 -203327 -108
```

```

## ..- attr(*, "names")= chr [1:178] "1" "2" "3" "4" ...
## $ effects      : Named num [1:178] -19740035 903638 1527237 1196025 1...
## ..- attr(*, "names")= chr [1:178] "(Intercept)" "advertising_spend"
## $ rank         : int 8
## $ fitted.values: Named num [1:178] 1101465 1177273 881460 1057583 103...
## ..- attr(*, "names")= chr [1:178] "1" "2" "3" "4" ...
## $ assign        : int [1:8] 0 1 2 3 4 5 6 7
## $ qr           :List of 5
##   ..$ qr     : num [1:178, 1:8] -13.342 0.075 0.075 0.075 0.075 ...
##   ... ..- attr(*, "dimnames")=List of 2
##   ... ..- attr(*, "assign")= int [1:8] 0 1 2 3 4 5 6 7
##   ..$ qraux: num [1:8] 1.07 1.15 1.13 1.13 1.07 ...
##   ..$ pivot: int [1:8] 1 2 3 4 5 6 7 8
##   ..$ tol   : num 1e-07
##   ..$ rank  : int 8
##   ..- attr(*, "class")= chr "qr"
## $ df.residual  : int 170
## $ xlevels      : Named list()
## $ call          : language lm(formula = model_2, data = sales_data)
## $ terms         :Classes 'terms', 'formula' language monthly_sales ~ ...
##   ... ..- attr(*, "variables")= language list(monthly_sales, advertising_spend)
##   ... ..- attr(*, "factors")= int [1:8, 1:7] 0 1 0 0 0 0 0 0 ...
##   ... ... ..- attr(*, "dimnames")=List of 2
##   ... ... ..- attr(*, "term.labels")= chr [1:7] "advertising_spend" "staff_size" ...
##   ... ... ..- attr(*, "order")= int [1:7] 1 1 1 1 1 1 1
##   ... ... ..- attr(*, "intercept")= int 1
##   ... ... ..- attr(*, "response")= int 1
##   ... ... ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
##   ... ... ..- attr(*, "predvars")= language list(monthly_sales, advertising_spend)
##   ... ... ..- attr(*, "dataClasses")= Named chr [1:8] "numeric" "numeric" ...

```

```

## .. . . . - attr(*, "names")= chr [1:8] "monthly_sales" "advertising_s
## $ model           :'data.frame':   178 obs. of  8 variables:
##   ..$ monthly_sales      : num [1:178] 678303 755785 831246 854256 92
##   ..$ advertising_spend : num [1:178] 71756 13050 60465 49036 12695
##   ..$ staff_count        : num [1:178] 11 16 5 16 24 16 14 13 22 19 .
##   ..$ store_size         : num [1:178] 224 1122 426 150 858 ...
##   ..$ foot_traffic       : num [1:178] 8477 6055 6837 5725 4243 ...
##   ..$ local_income        : num [1:178] 12000 17894 17044 29601 12000
##   ..$ competition_distance: num [1:178] 0.4 5.4 3.3 6 6.8 2.9 0.4 8.8 3
##   ..$ seasonal_index      : num [1:178] 0.987 1.156 0.7 0.978 0.989 ...
##   ..- attr(*, "terms")=Classes 'terms', 'formula' language monthly_sa
##   .. . . . - attr(*, "variables")= language list(monthly_sales, advertisi
##   .. . . . - attr(*, "factors")= int [1:8, 1:7] 0 1 0 0 0 0 0 0 0 0 ...
##   .. . . . . - attr(*, "dimnames")=List of 2
##     .. . . . - attr(*, "term.labels")= chr [1:7] "advertising_spend" "sta
##     .. . . . - attr(*, "order")= int [1:7] 1 1 1 1 1 1 1
##     .. . . . - attr(*, "intercept")= int 1
##     .. . . . - attr(*, "response")= int 1
##     .. . . . - attr(*, ".Environment")=<environment: R_GlobalEnv>
##     .. . . . - attr(*, "predvars")= language list(monthly_sales, advertisi
##     .. . . . - attr(*, "dataClasses")= Named chr [1:8] "numeric" "numeric"
##     .. . . . . - attr(*, "names")= chr [1:8] "monthly_sales" "advertisin
##   - attr(*, "class")= chr "lm"

```

Task 5: Interpret your final model results in business terms. What do the coefficients tell us about the drivers of sales performance?

Business Interpretation:

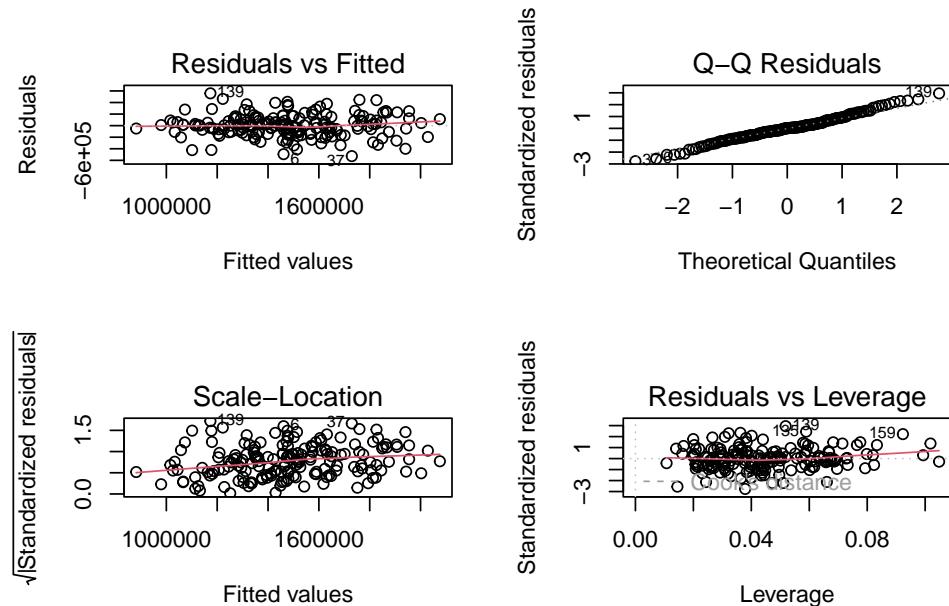
The business should lean more into Investments in marketing, staffing, and strategies to increase store visits can positively impact sales and it appears that Local economic conditions also matter in the consideration of locations since stores in

higher-income areas tend to perform better.

21 Validation and Diagnostics

Check whether your selected model meets the assumptions of linear regression.

```
# Your code here: Create diagnostic plots  
# Check residual patterns, normality, and influential observations  
  
# Create diagnostic plots  
par(mfrow = c(2, 2))  
plot(final_fit)
```



```
par(mfrow = c(1, 1))  
  
# Additional diagnostic tests  
residuals_data <- tibble(  
  fitted = fitted(final_fit),  
  residuals = residuals(final_fit),
```

```

standardised_residuals = rstandard(final_fit)
)

# Calculate diagnostic statistics
leverage_threshold <- 2 * length(coef(final_fit)) / nrow(sales_data)
high_leverage <- sum(hatvalues(final_fit) > leverage_threshold)
influential_points <- sum(abs(rstudent(final_fit)) > 2)

```

Task 6: Comment on model assumptions and any potential concerns with your final specification.

Model Diagnostics Assessment:

The final model explains about 43% of monthly sales variation (Adjusted $R^2 = 0.421$). Residuals are roughly normal with mostly linear relationships, though variability slightly increases at higher predicted sales. No extreme outliers were detected, but predictions at very high or low sales should be interpreted cautiously

22 Business Recommendations

Based on your analysis, provide actionable recommendations for RetailMax management.

Task 7: Create specific, data-driven recommendations for improving sales performance. Consider both the statistical evidence and practical business constraints.

Management Recommendations:

1. **Priority Actions:** *Increase advertising spend strategically, Optimize staffing levels, Leverage foot traffic and local income data*
2. **Resource Allocation:** *Prioritize investment in advertising channels that historically drive measurable sales increases and monitor underperforming stores for potential adjustments in staffing or local marketing strategies.*

3. **Performance Monitoring:** *Track monthly sales, advertising ROI, staff efficiency metrics, and foot traffic trends.*
4. **Model Limitations:** *The model explains 43% of sales variance, leaving a substantial portion unexplained. Factors like seasonality, competitor activity, or online sales may also influence sales. Future analyses could incorporate additional data, such as promotions, competitor pricing, or customer demographics, to improve predictive accuracy.*

23 Conclusion

Cross-validation provides essential protection against overfitting when comparing different model specifications. Your analysis shows how systematic model comparison enables evidence-based decision-making in business contexts.

The step-by-step process of model development, validation, and interpretation ensures that analytical insights translate effectively into actionable business strategy whilst maintaining statistical rigour.

Task 8: Reflect on the cross-validation process. How did the CV results influence your model selection compared to simply examining in-sample fit statistics?

Cross-Validation Reflection:

Cross-validation provided a realistic estimate of model performance on new data. It showed that simpler models performed nearly as well as the full model, helping avoid overfitting. This confirmed that the chosen model balances predictive accuracy with interpretability and is likely to generalize well.